

Cardinality-Constrained Portfolio Optimization

A Genetic Algorithm Out-of-Sample Evaluation

Georgios Dedempilis (2803764)
BSc Computer Science - VU Amsterdam
Supervisor: Prof. dr. A.E. Eiben
June 2026

Table of contents

01

Motivation &
Research Questions

02

Algorithm &
Experimental Setup

03

Results

04

Conclusions &
Future Work

05

Threats to Validity

01

Motivation & Research Questions

Motivation



THE PROBLEM

- Real portfolios cannot hold 500+ stocks efficiently.
- Transaction costs, complexity, and regulation matter.
- Markowitz MVO ignores these practical constraints.



THE GAP

- Cardinality constraints make the problem NP-hard.
- Standard QP solvers cannot handle it, metaheuristics are needed.
- GAs are a natural fit, but out-of-sample evidence remains limited.

Can a GA beat simpler benchmarks across 20 years of real market data, and if not, why not?

Research Questions



RQ1 What is the impact of a 0.15 maximum weight bound on out-of-sample risk-adjusted performance vs unconstrained MVO?



RQ2 Do evolutionary algorithms achieve competitive Sharpe ratio vs constrained MVO, unconstrained MVO, and 1/N under realistic constraints?



RQ3 How does the choice of cardinality K affect GA Sharpe ratio, turnover, and concentration, and how does adaptive K compare to fixed- K configurations?

02

Algorithm & Experimental Setup

Algorithm



Why a Genetic Algorithm?

NP-hard problem - QP solvers fail. Real-valued chromosome fits portfolio weights better than ES or DE. Chang (2000), Streichert (2003).



Chromosome Representation

$N \approx 867$ positions, $K \in [10, 30]$ non-zero entries each $\in [0.02, 0.15]$, $\sum w = 1$

0	0.08	0	0.15	0.05	0	0.10
---	------	---	------	------	---	------

...

0	0.09	0.10	0.02	0	0.04	0
---	------	------	------	---	------	---



Live Walkthrough

Steps: population -> selection -> crossover -> mutation -> repair + local refinement -> new generation (elitism + median-fitness run selection)

Experimental Setup

Data

- CRSP US equities, Jan 2000 - Dec 2025
- Approximately 867 eligible stocks/month (NYSE + NASDAQ, market cap $\geq$ \$2B)
- 252 out-of-sample rebalancing periods (Jan 2005 - Dec 2025)
- Risk-free rate: FRED 3-month T-bill (DTB3)



Setup

Estimation window: 60 months rolling

Covariance: Ledoit-Wolf shrinkage

Optuna tuning: 2005-2012, 15 Trials, TPE

Experimental Setup

Benchmarks



Constrained MVO: weights $\leq 15\%$



Unconstrained MVO: weights $\leq 100\%$



1/N: equal weights for ~867 stocks

Costs




Transaction cost $\gamma = 0.3\%$ on all models

Model	Weight bounds	Cardinality	Estimation
GA	[0.02, 0.15]	$K \in [10, 30]$	Yes
Constrained MVO	[0.00, 0.15]	None	Yes
Unconstrained MVO	[0.00, 1.00]	None	Yes
1/N	1/N each	All ~867	No

03

Results

Cumulative returns



All returns net of transaction costs ($\gamma = 0.3\%$). Initial portfolio value normalised to 1.0. Shaded regions indicate the Global Financial Crisis (Sep 2008 – Jun 2009) and COVID-19 shock (Feb – Apr 2020).

By December 2025 both MVO variants and 1/N reach **higher** terminal value than the GA.

Main performance table

Metric	GA	Const. MVO	Unconst. MVO	1/N
Sharpe (net)	0.2741	0.5810	0.5809	0.5307
Ann. Return	5.46%	8.28%	8.28%	9.13%
Ann. Vol.	19.92%	14.26%	14.26%	17.20%
Max DD	-55.76%	-43.07%	-43.07%	-51.31%
Turnover	22.19%	17.07%	17.07%	4.50%

The GA **underperforms** all three benchmarks. Both MVO variants reach nearly identical Sharpe ratios, the 0.15 weight cap barely binds across 867 stocks.

Statistical significance

Paired t-test (mean return difference)
Jobson-Korkie test with Memmel correction (Sharpe ratio difference)

Comparison	Sharpe Diff	JK p-value	JK Sig.	t p-value	t Sig.
GA vs 1/N	-0.257	0.1316	no	0.2558	no
GA vs Unconst. MVO	-0.307	0.0045	yes	0.2051	no
GA vs Const. MVO	-0.307	0.0045	yes	0.2051	no

The Sharpe underperformance against both MVO variants is **statistically significant** (JK, $p < 0.01$), but the mean return difference is not significant for any benchmark.

The central finding: λ -ablation

Full 252-period re-run with penalty removed ($\lambda = 0$) vs tuned ($\lambda = 1.8437$).


	Tuned $\lambda = 1.8437$	$\lambda = 0$ (no penalty)
Net Sharpe	0.2741	0.4993
Avg cardinality K	16.3	25.1
Avg turnover	22.2%	57.1%

Removing the penalty nearly doubles Sharpe, the penalty suppresses cardinality too **aggressively**, costing more in return than it saves in trading costs.

Fixed-K sensitivity

Fixed hyperparameters held constant, only K varies across {10, 15, 20, 25, 30}.

Figure K1: Effect of Fixed Cardinality K on Performance Metrics



Sharpe rises with K up to K = 25-30, while concentration falls monotonically. The adaptive GA settles at K = 16.3, below every fixed-K configuration except K = 10.

04

Conclusions & Future Work

Conclusions

RQ1 (constraints vs unconstrained MVO): Weight constraints have negligible impact, constrained and unconstrained MVO reach near-identical Sharpe (0.5810 vs 0.5809) across 867 stocks.

RQ2 (GA vs benchmarks): GA underperforms all benchmarks (Sharpe 0.2741 vs 0.53–0.58, JK $p=0.004$ vs MVO). Root cause: over-aggressive turnover penalty, removing it raises Sharpe to 0.4993.

RQ3 (cardinality K): Higher K improves Sharpe (peaks K=25–30), concentration falls monotonically. Turnover penalty suppresses adaptive K to 16.3, below all fixed-K configs except K=10.

Key takeaway: The main bottleneck is the turnover penalty, not the evolutionary approach.

Future Work



Extend to international
or multi-asset portfolios



Scale to 30 parallel GA
runners



Walk-forward Optuna
tuning



Tune λ per fixed-K
configuration

05

Threats to Validity

Threats to Validity



Internal

8 runners instead of 30, relaxed MVO solver tolerances, 105 tests validate implementation.

External

US large/mid-cap only, GFC and COVID period not representative of average conditions.

Construct

Sharpe assumes normality, tuning window overlaps evaluation (mild data-snooping).

Conclusion

JK rejects GA vs MVO ($p=0.0045$), paired t-test does not ($p=0.205$), tests answer different questions.

Thank you!

Do you have any questions?

g.dedempilis@student.vu.nl

Georgios (George) Dedempilis

Student nr: 2803764

VunetID: gde218

github.com/georgeded/ga-portfolio-optimization

